

# SIA - TP5

---

Grupo 7

Tomás Scheffer - 63393  
Matías Sapino - 61067  
Tobías Pugliano - 62180  
Luca Bloise - 63004

# Autoencoder

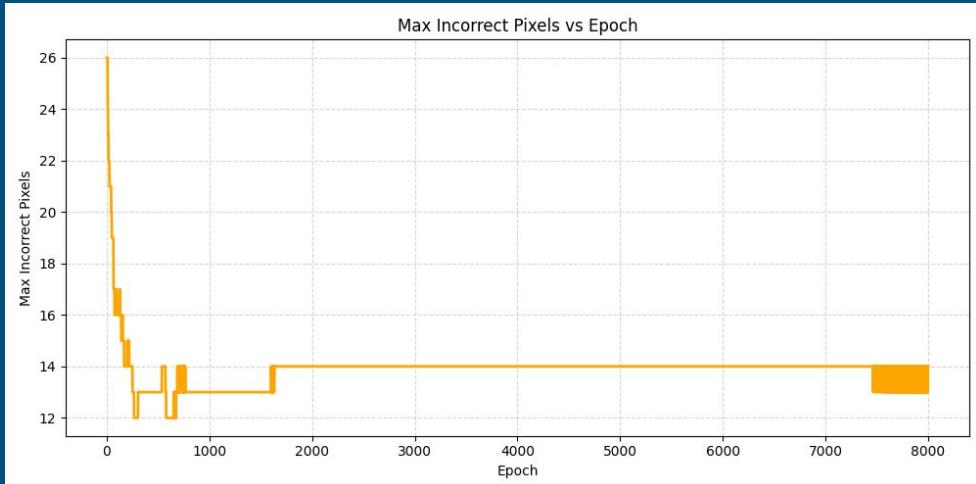
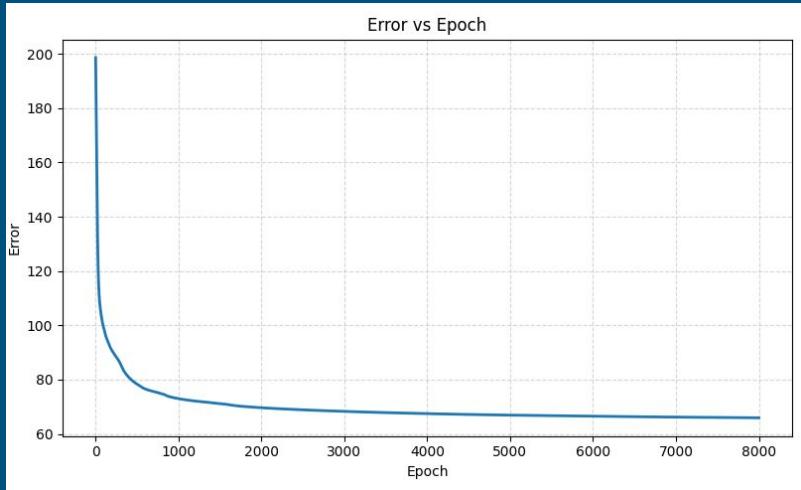
---

# Arquitectura de red e hiperparametros

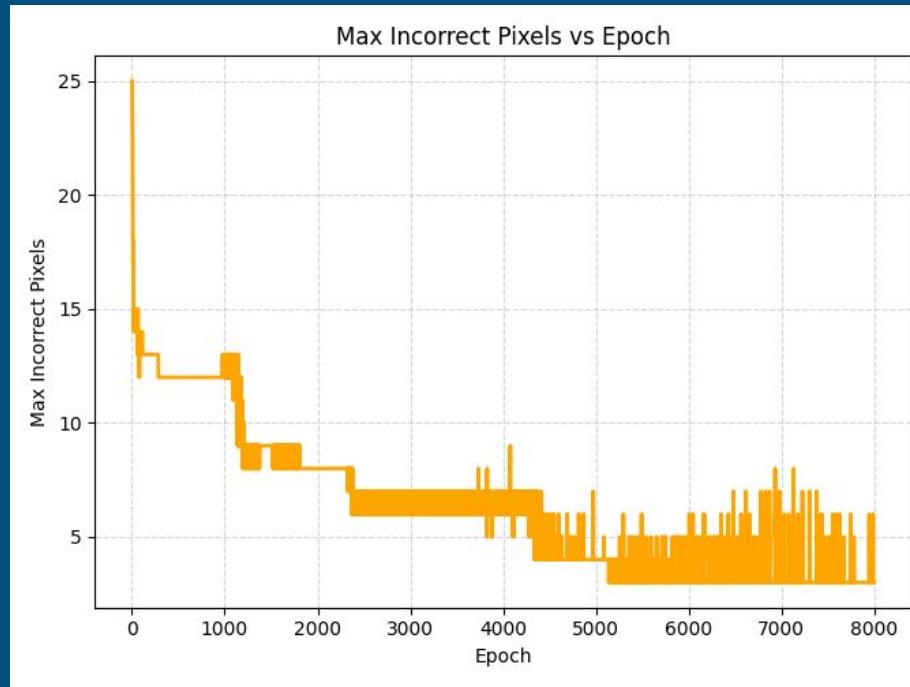
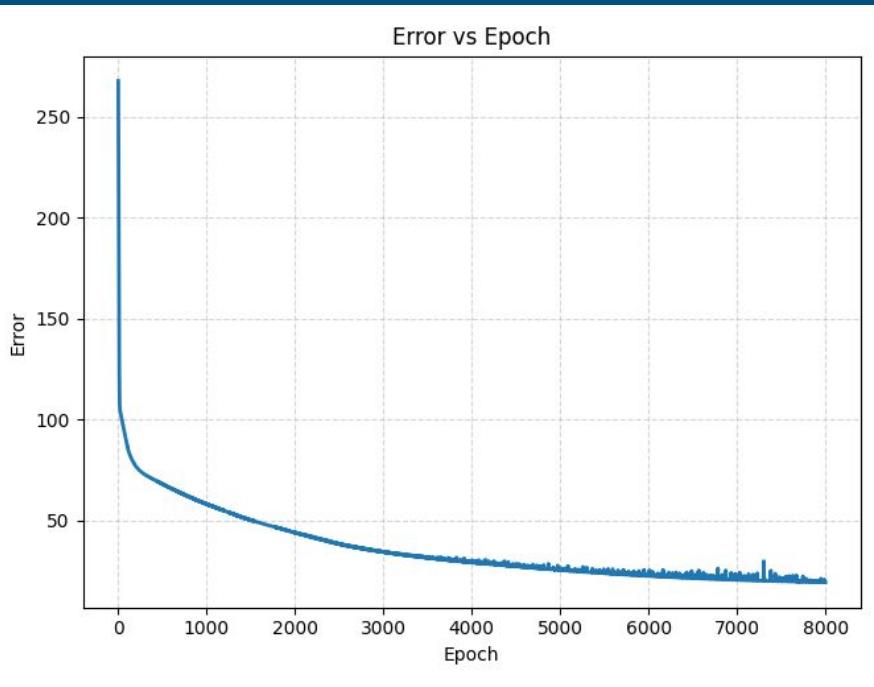
---

- Tomamos como base la implementación del perceptrón multicapa realizada en el TP3
- En dicho TP se obtuvo que el mejor optimizador para el perceptrón es Adam por lo cual optamos por utilizar el mismo
- Comenzamos con una arquitectura pequeña y vamos aumentando la cantidad de capas intermedias del encoder y decoder, ya que en el TP3 notamos que a mayor cantidad de capas intermedias menor resultaba el error
- Las siguientes pruebas se realizaron con un learning rate  $\eta = 0.000985$

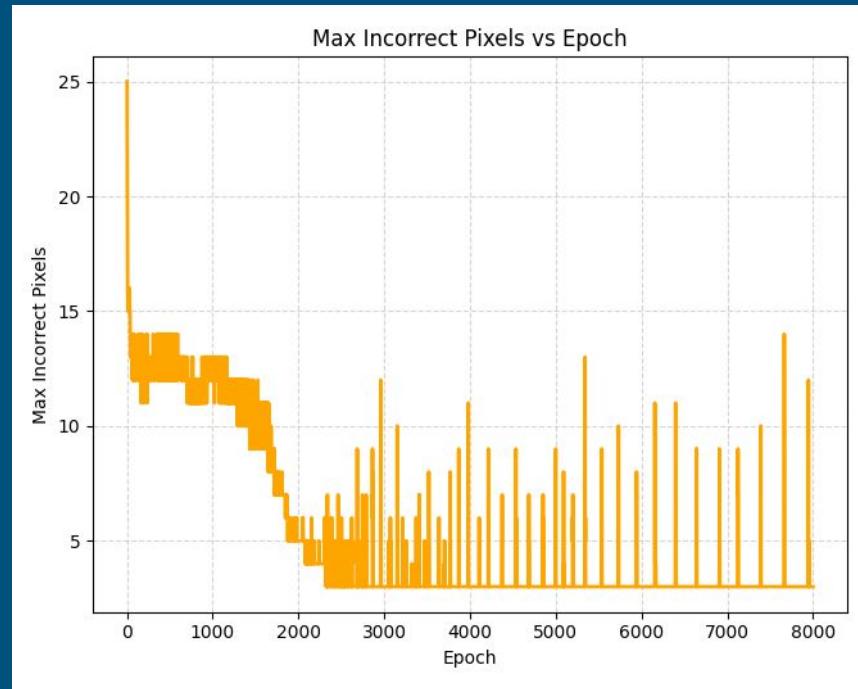
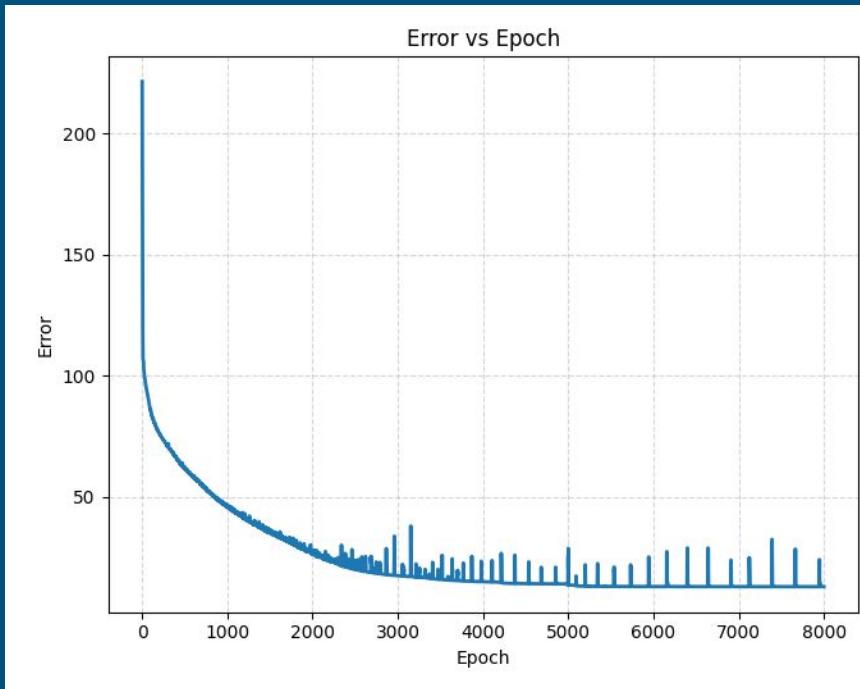
# Modificando la estructura: [35, 2, 35]



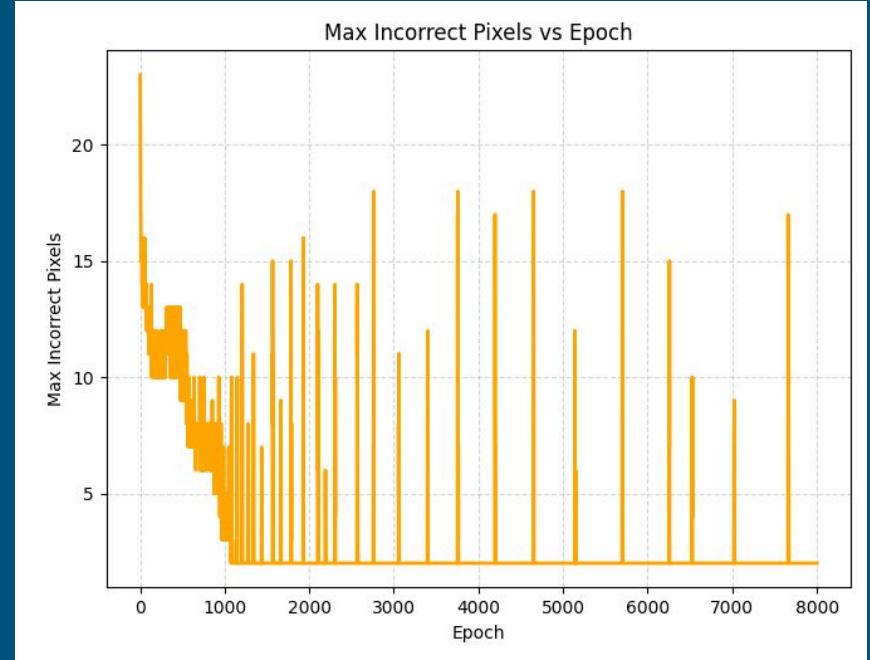
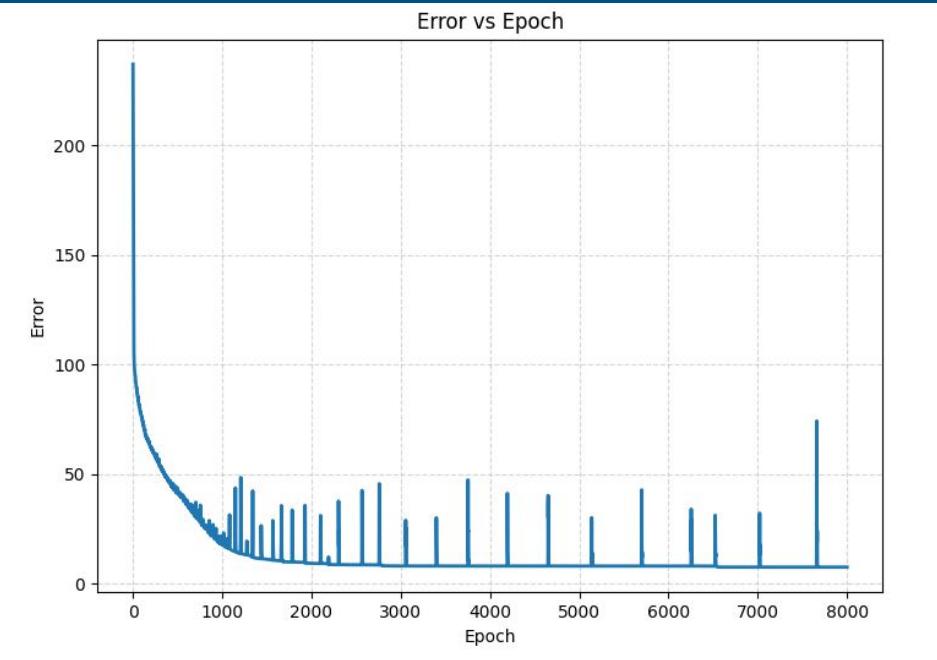
# Modificando la estructura: [35, 16, 2, 16, 35]



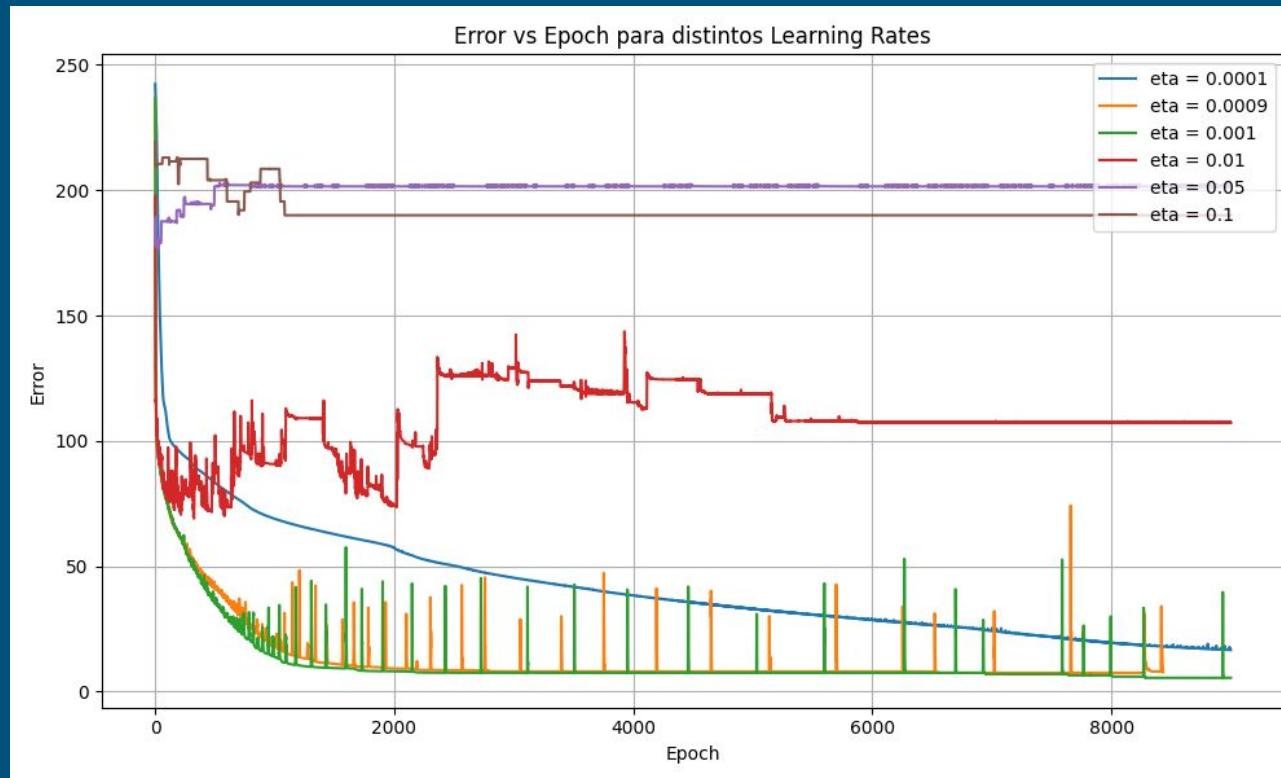
# Modificando la estructura: [35, 24, 16, 2, 16, 24, 35]



# Modificando la estructura: [35, 32, 24, 16, 2, 16, 24, 32, 35]



# Modificando el learning rate

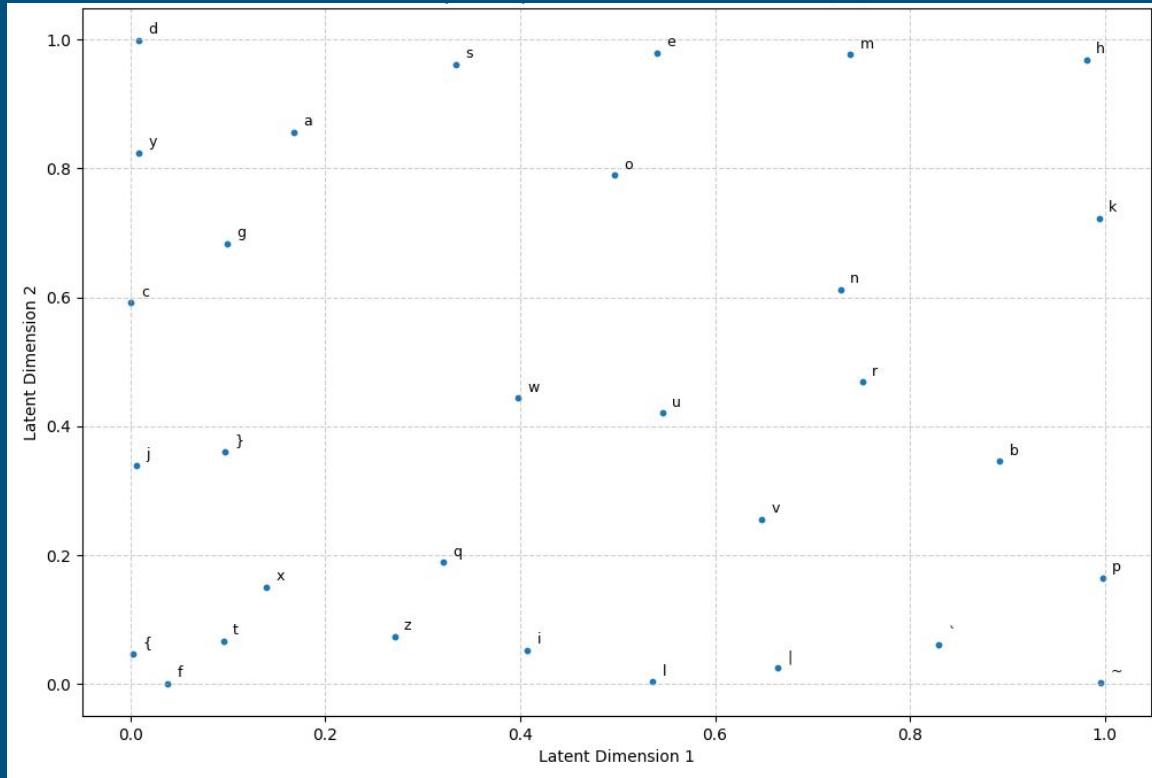


# Arquitectura de red e hiperparametros

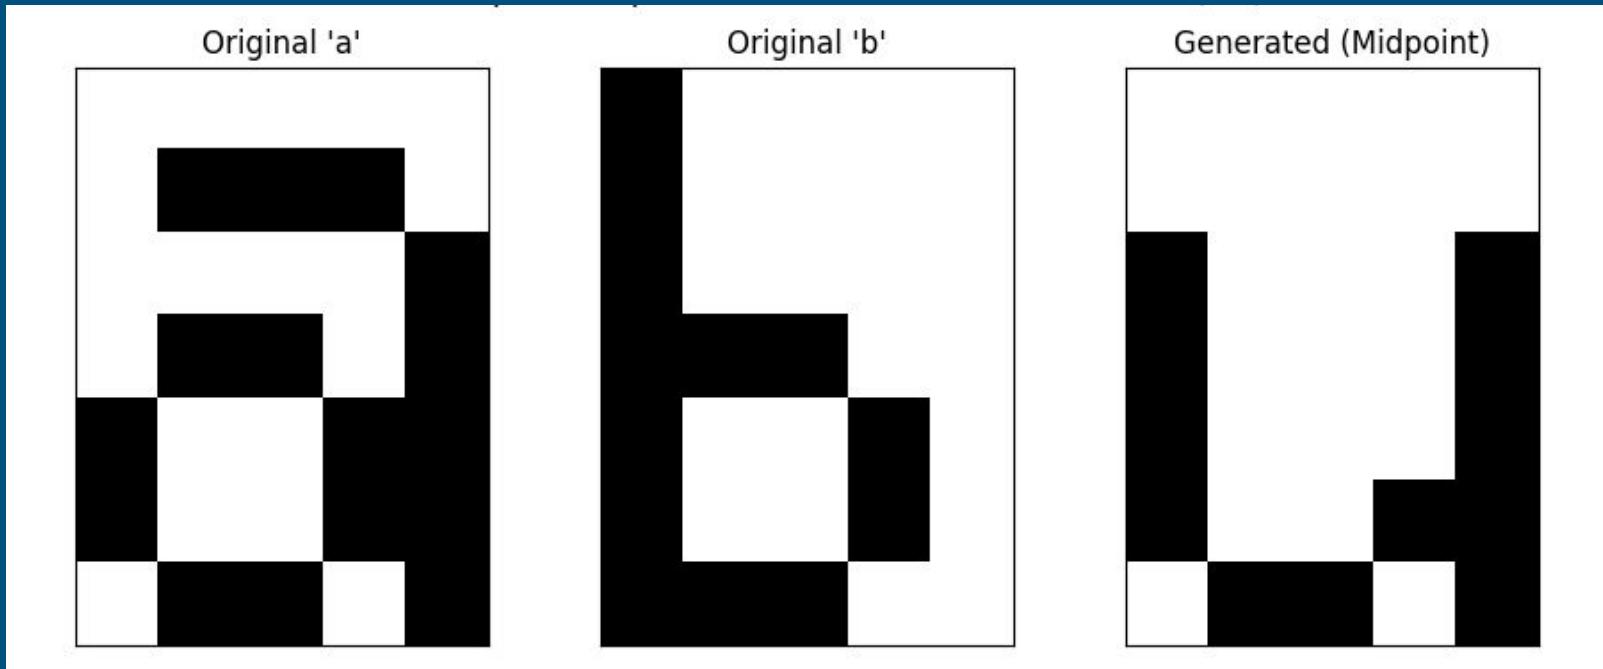
---

- Encoder y Decoder de 4 capas conectados por un espacio latente de 2 neuronas ( [35, 32, 24, 16, 2, 16, 24, 32, 35] )
- Como función de activación se utilizó la sigmoide
- Como learning rate se optó por  $\eta = 0.0009$

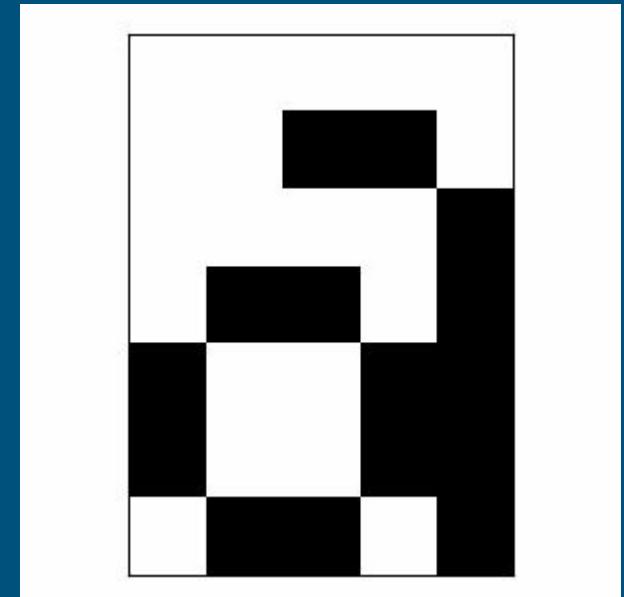
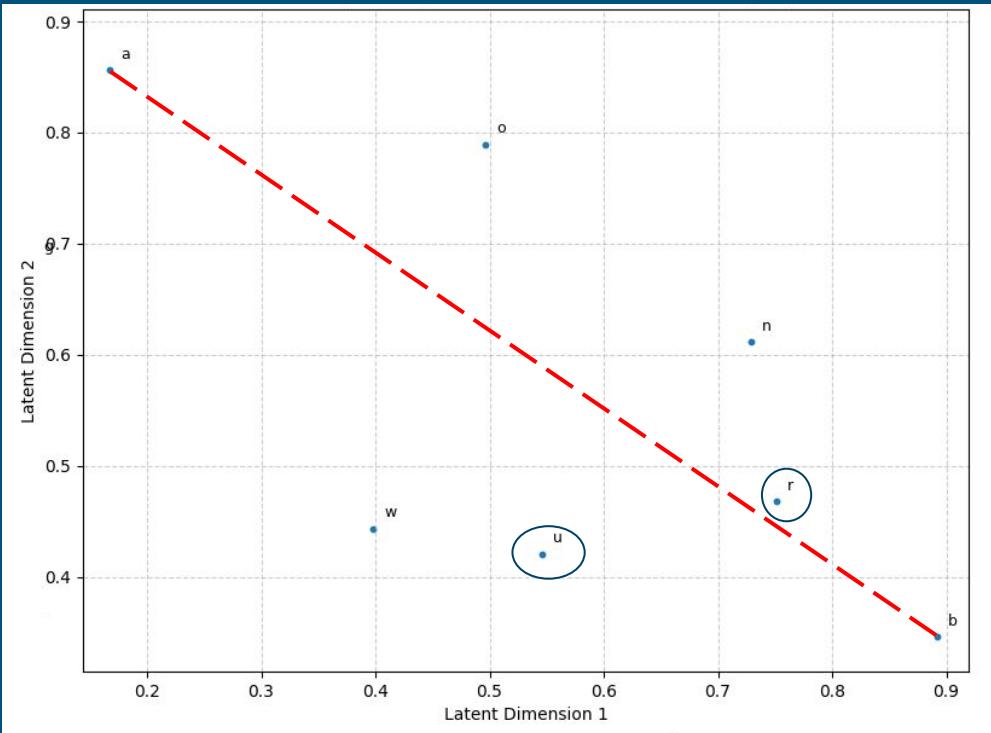
# Espacio Latente resultante



# Generación de una nueva letra



# Generación de una nueva letra



# Denoising Autoencoder

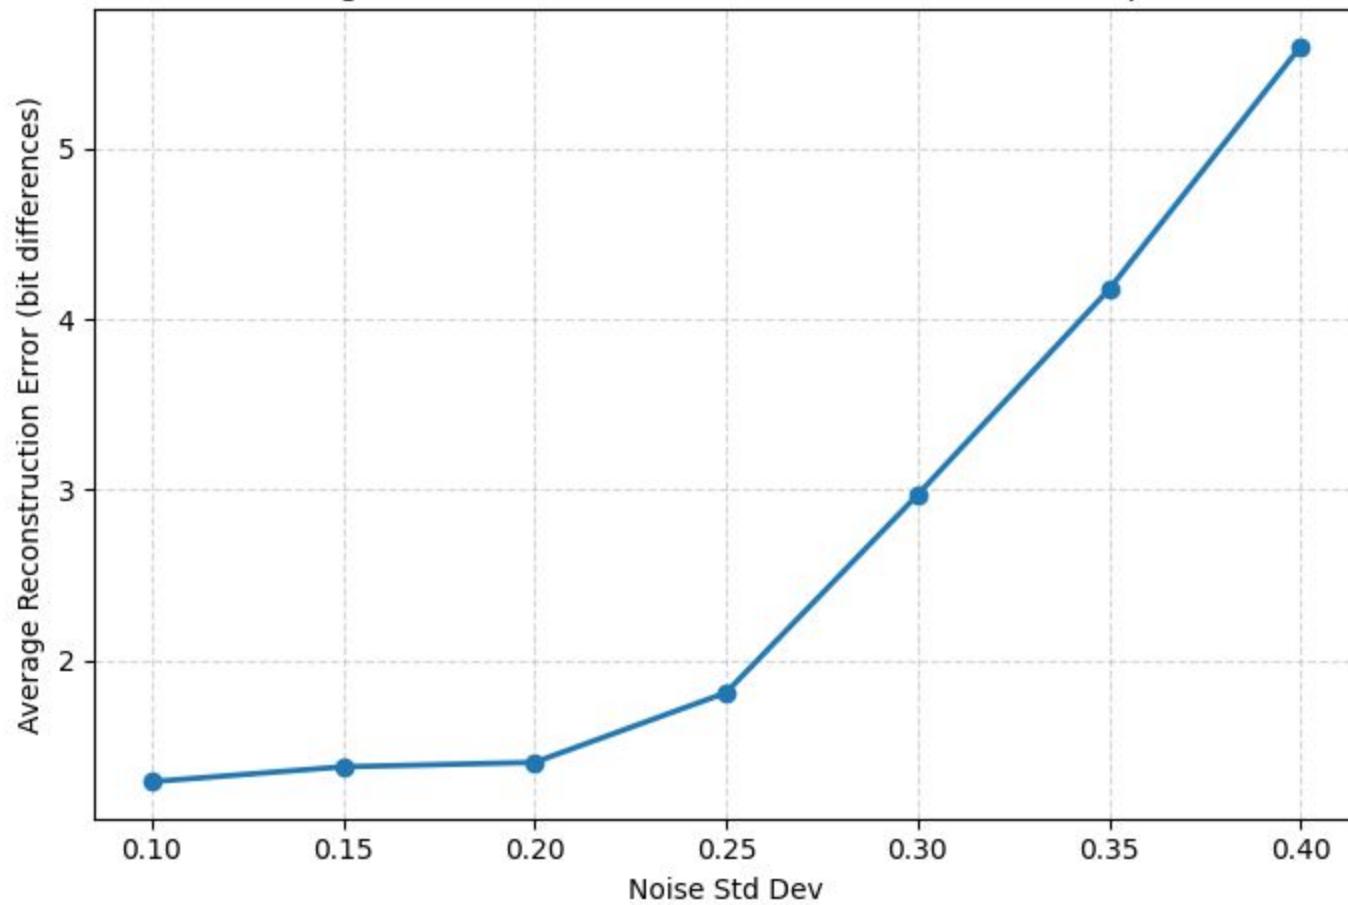
---

# Configuración y entrenamiento

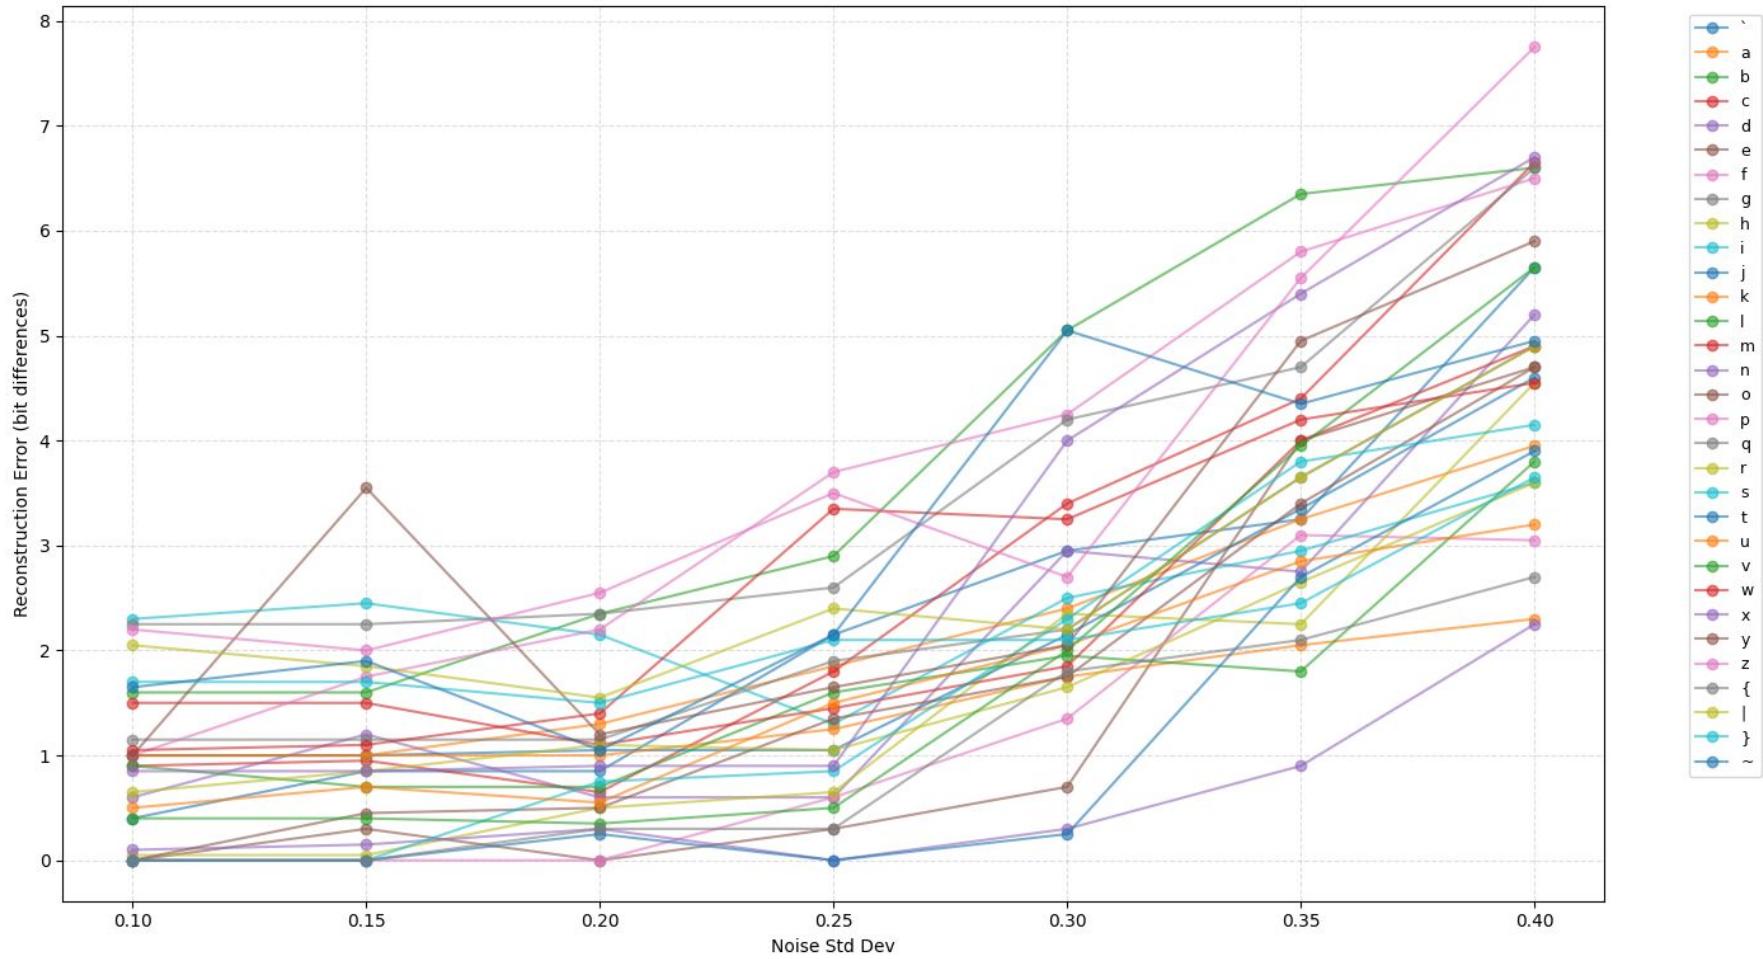
---

- Encoder y Decoder de 4 capas conectados por un espacio latente de 2 neuronas ( [35, 32, 24, 16, 2, 16, 24, 32, 35] )
- Función de activación sigmoide, learning rate  $\eta = 0.000985$ , 2000 Epochs
- Requirió varias pruebas encontrar cómo entrenar la red para conseguir resultados consistentes: usamos ruido binarizado y 5 versiones de cada carácter (stddev de 0.2)

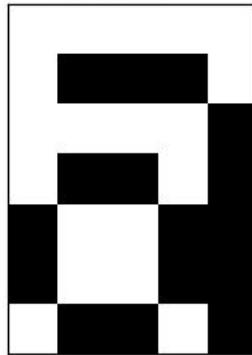
Average Reconstruction Error vs Noise Level (20 samples)



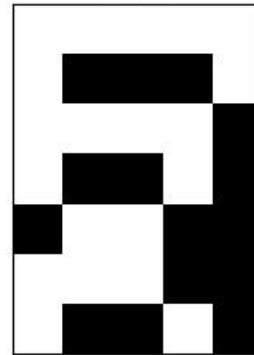
### Per-Character Reconstruction Error vs Noise Level (20 runs)



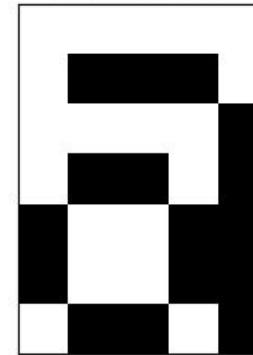
Original 'a'



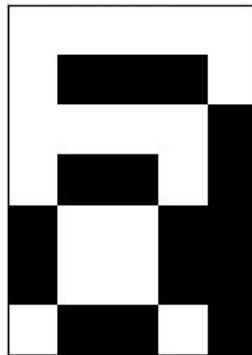
Noisy



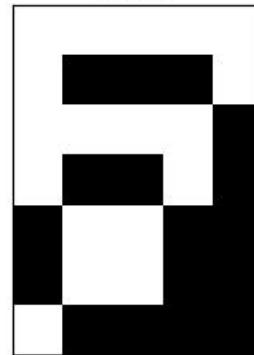
Reconstructed



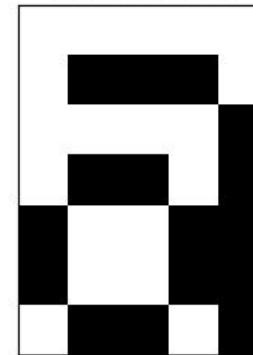
Original 'a'

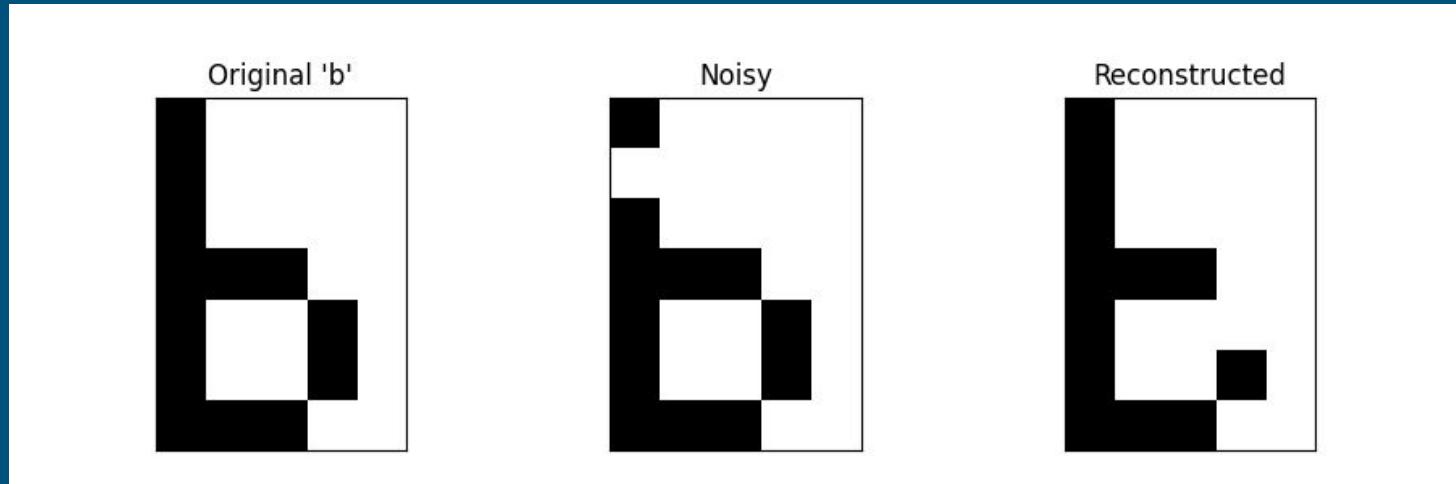
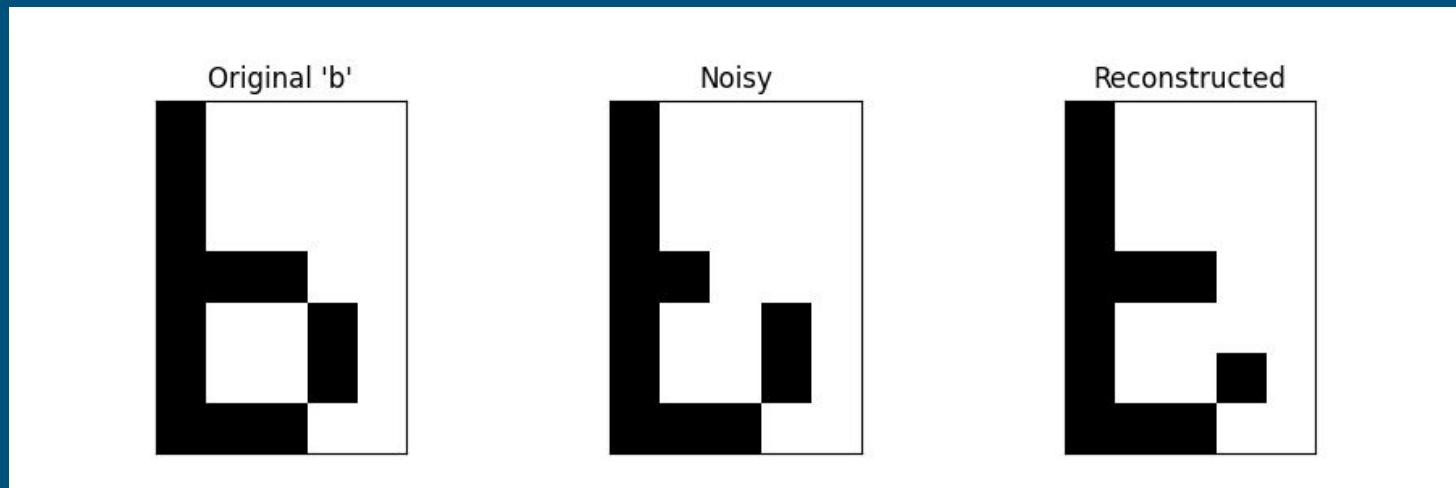


Noisy

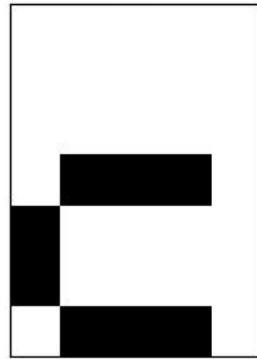


Reconstructed

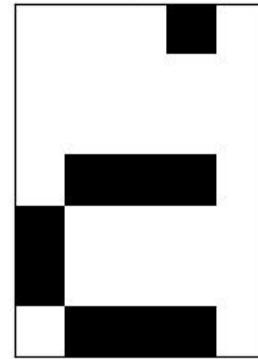




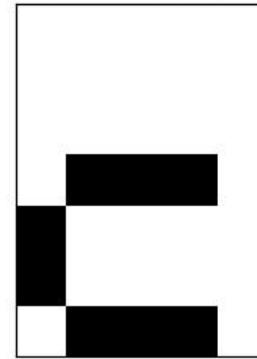
Original 'c'



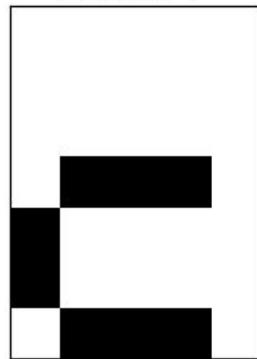
Noisy



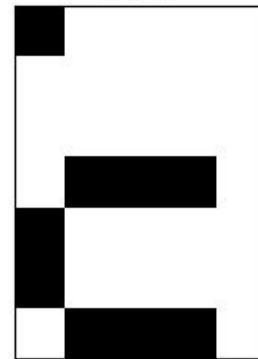
Reconstructed



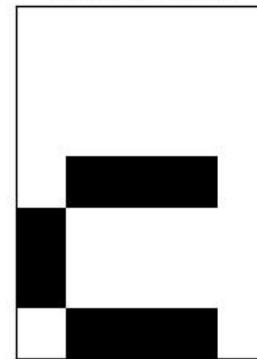
Original 'c'

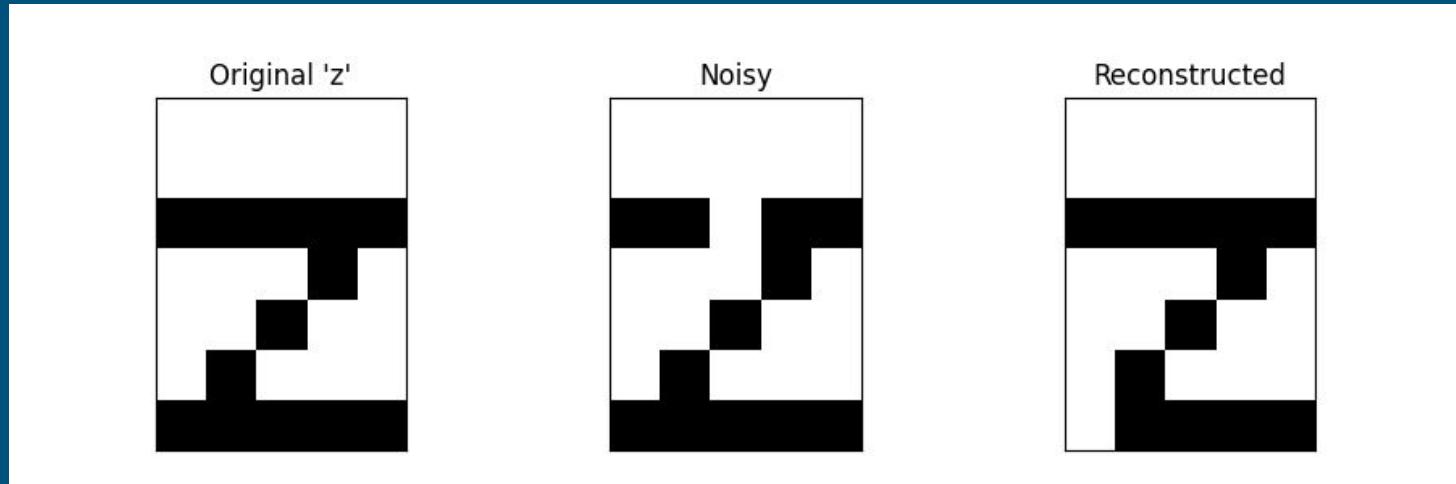
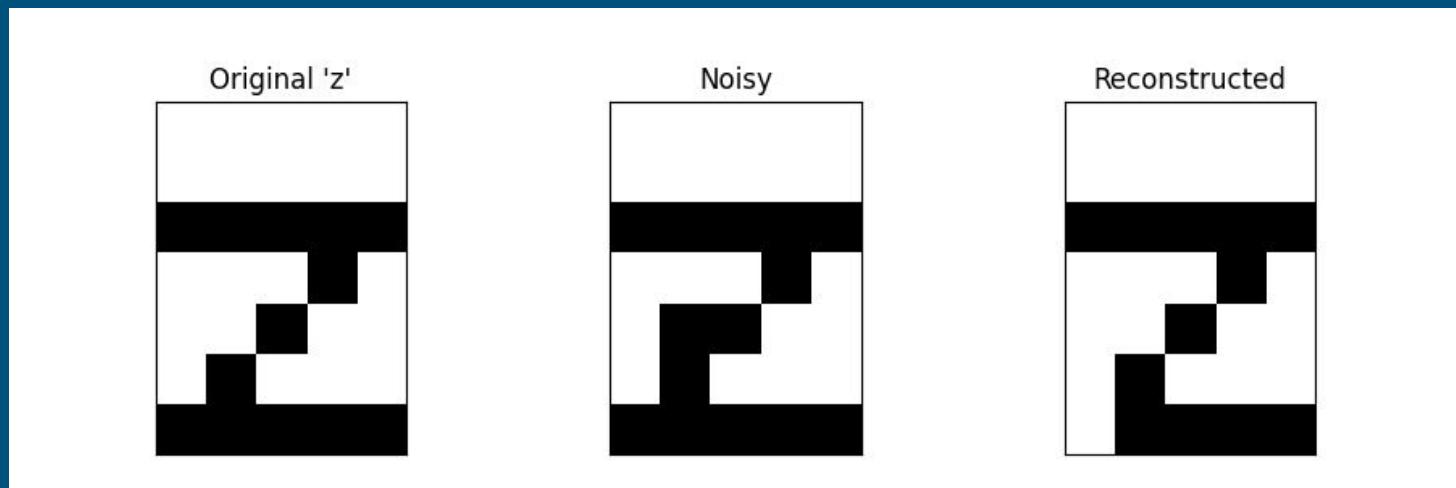


Noisy



Reconstructed



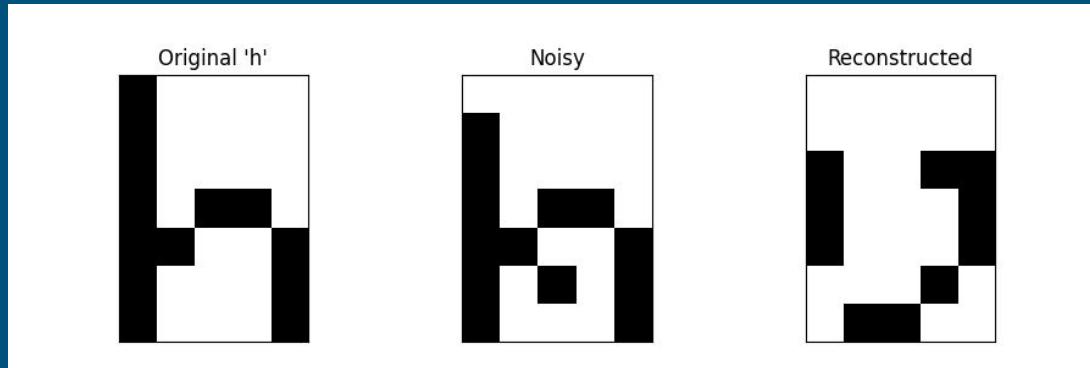
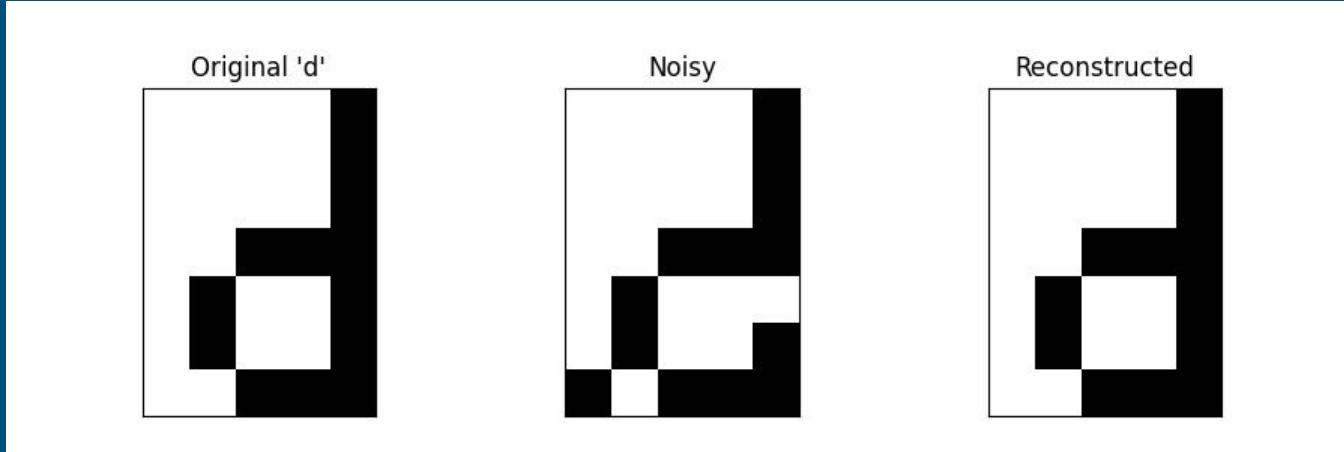


# Casos con más ruido

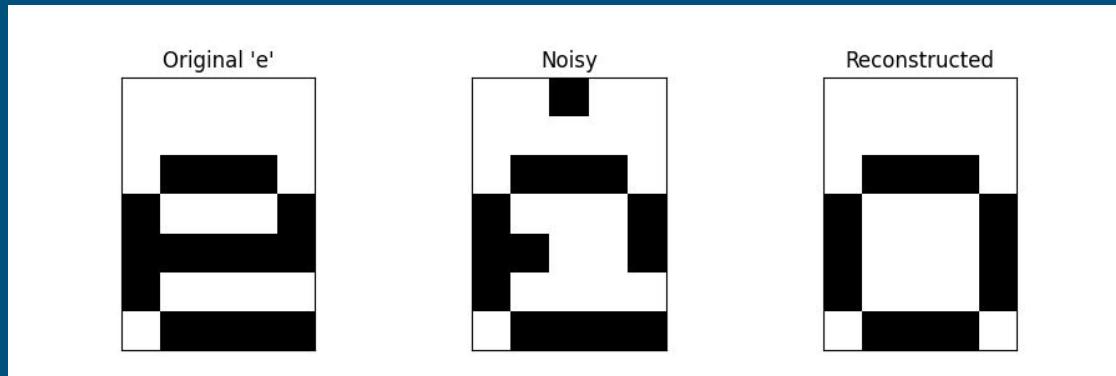
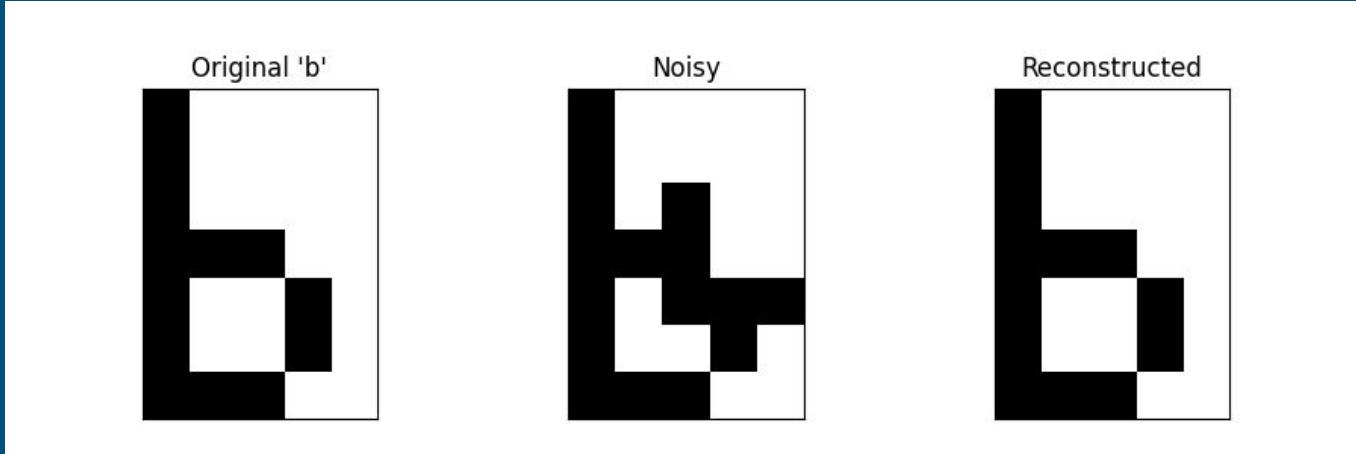
---

- Con más de un bit de diferencia es más propenso a fallar catastróficamente en la reconstrucción para ciertos caracteres
- Mientras que otros caracteres resultaron más “inmunes” al ruido
- Como se vió en el gráfico de Errores vs Noise Level por caracteres obtuvimos una distribución amplia

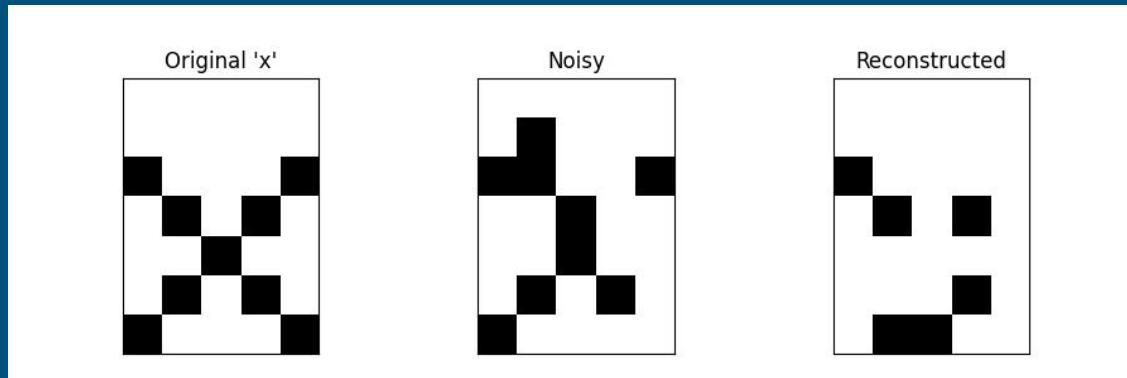
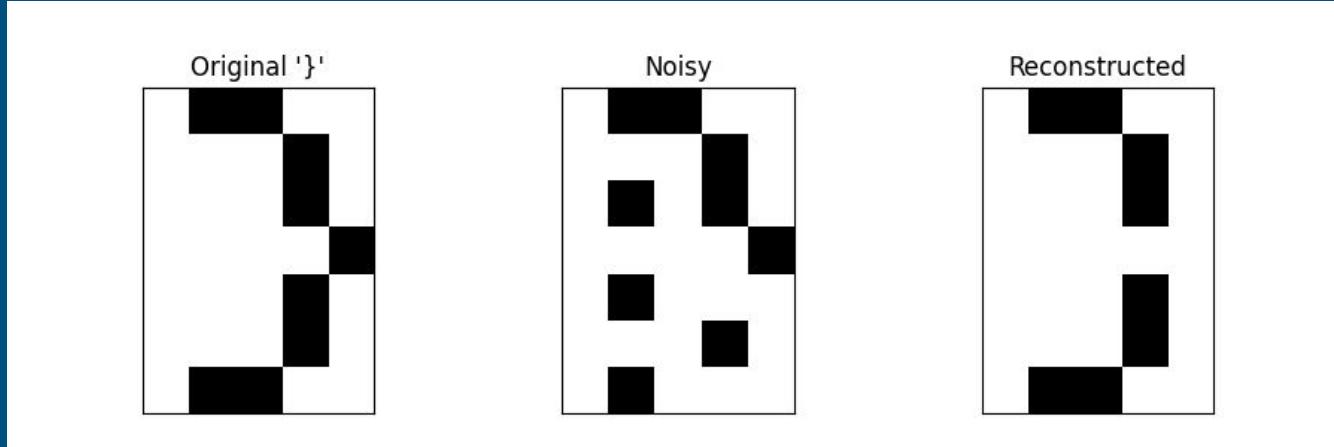
# 2 Flips



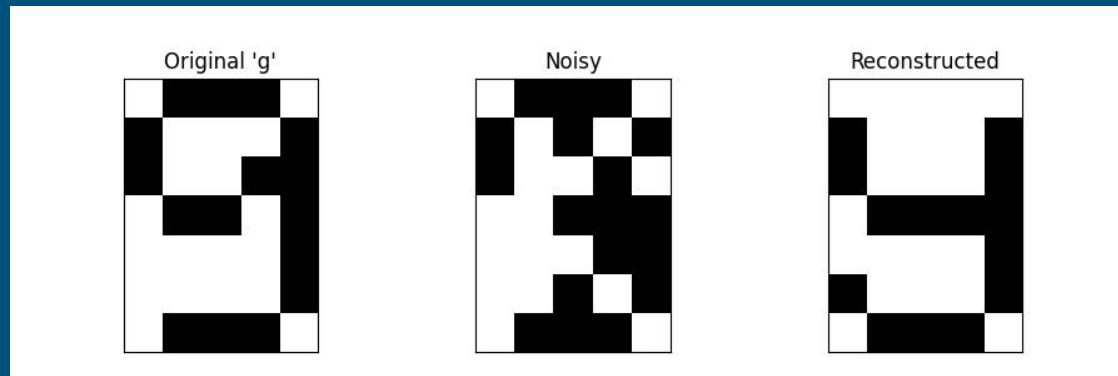
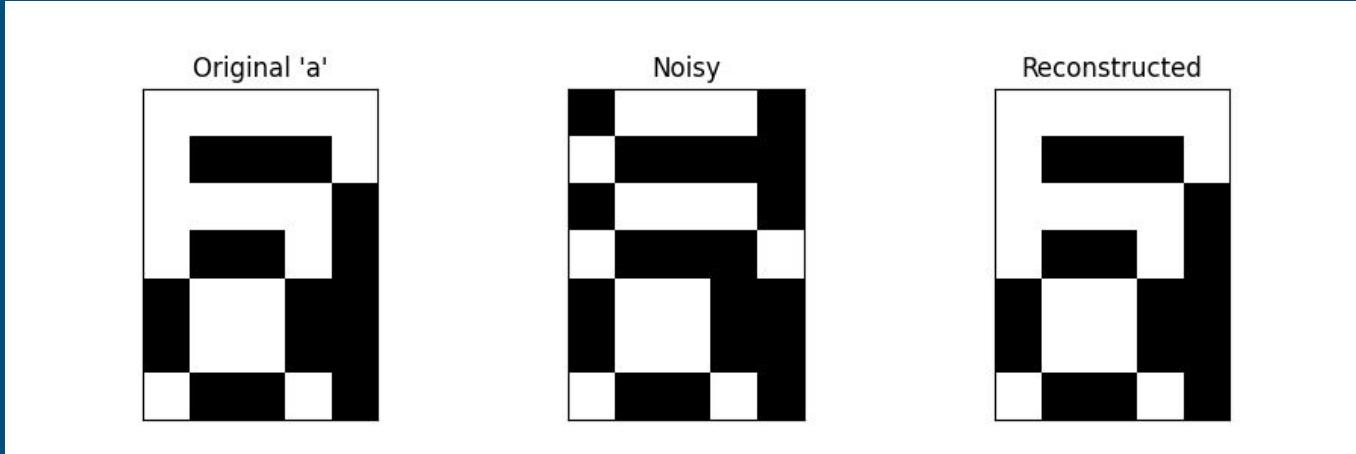
# 3 Flips



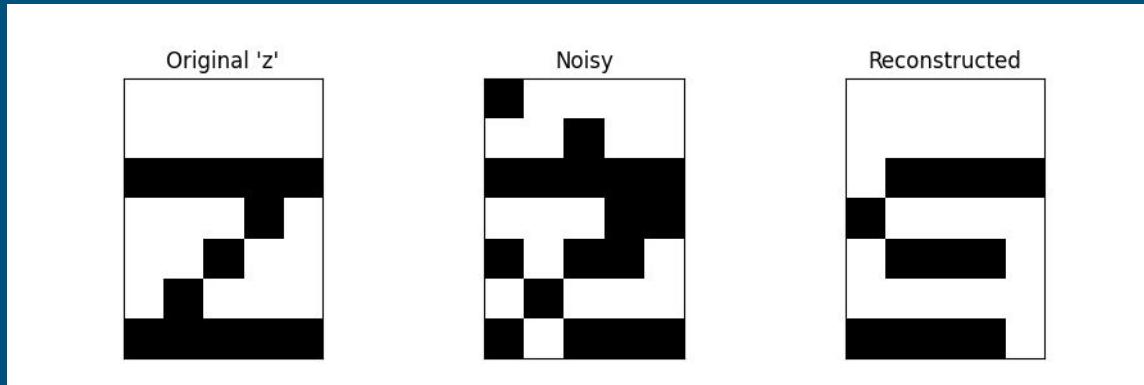
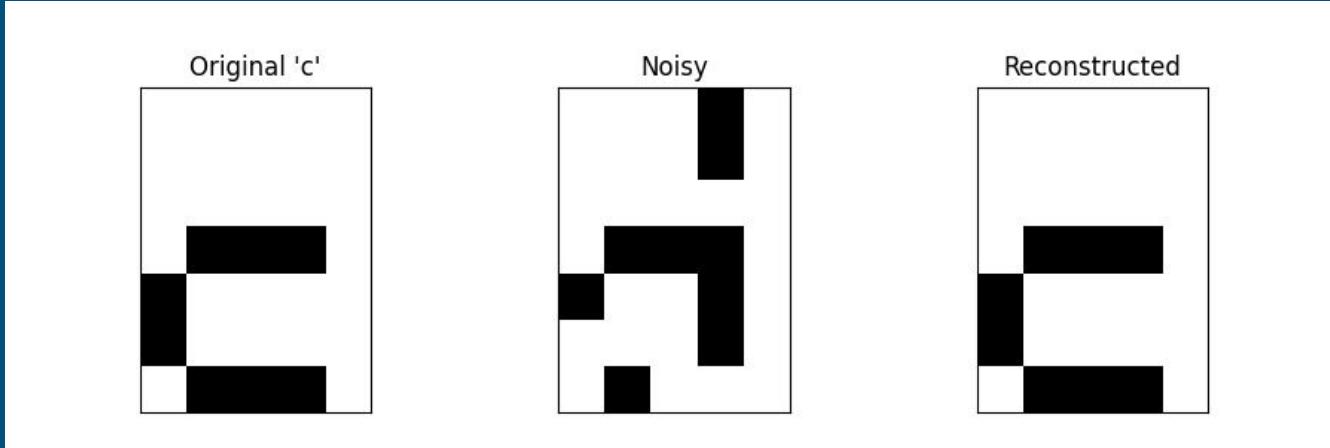
# 4 Flips



# 5 Flips



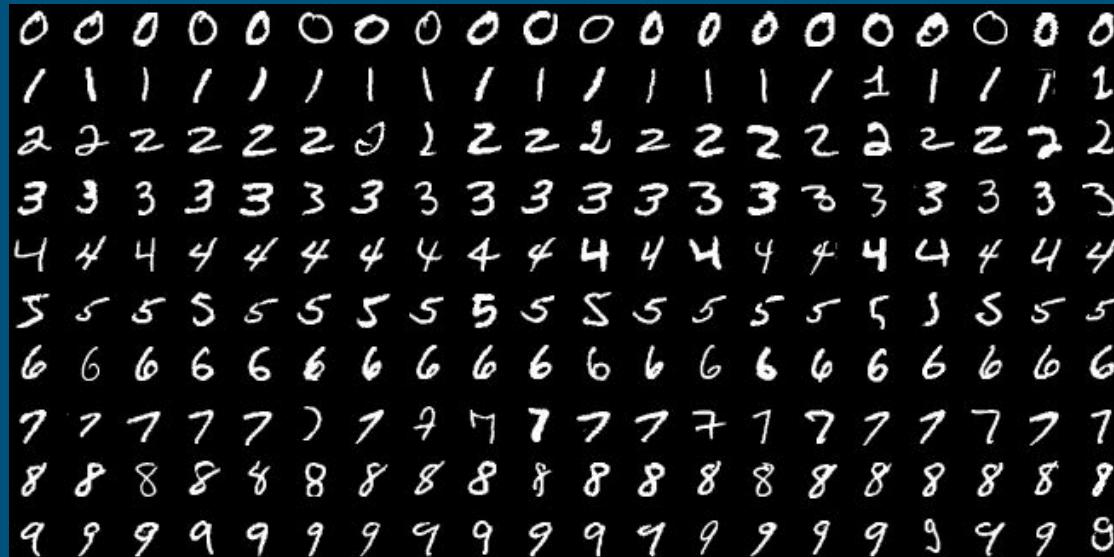
# 6 Flips



# Autoencoder Variacional

---

# Conjunto de datos para entrenar



20K elementos del conjunto de datos de MNIST (28x28)

# Arquitectura de red e hiperparametros

---

- Arquitectura:
    - Encoder: 784 -> 128 -> 64
    - Espacio Latente: 2
    - Decoder: 64 -> 128 -> 784
  - Optimizador: ADAM
  - Learning Rate: 0.001
- 
- Funciones de activación:
    - ReLU (Capas ocultas)
    - Sigmoidea (Última capa)
  - Epochs: 200
  - Batch Size: 64 (Mini-batch)

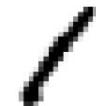
# Reconstrucción ( $E = 0.175$ )

VAE Reconstruction: Original vs Reconstructed (One per Digit)

Original 0



Original 1



Original 2



Original 3



Original 4



Original 5



Original 6



Original 7



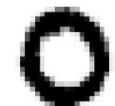
Original 8



Original 9



Reconstructed



Reconstructed



Reconstructed



Reconstructed



Reconstructed



Reconstructed



Reconstructed



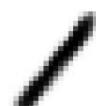
Reconstructed



Reconstructed



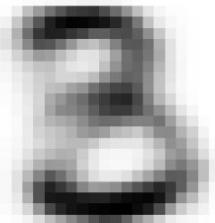
Reconstructed



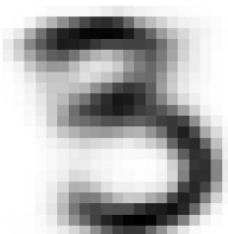
# Generación ( $E = 0.175$ )

VAE Random Generation with Latent Coordinates

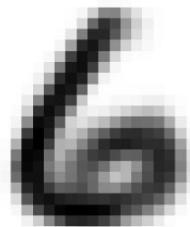
(-0.02, 0.46)



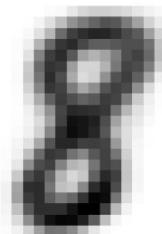
(0.16, 0.30)



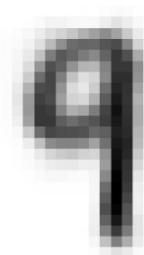
(-2.02, 1.02)



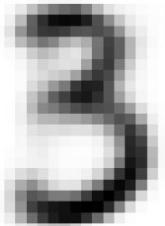
(-0.63, -0.39)



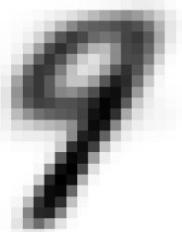
(1.33, -0.14)



(-0.79, 0.01)



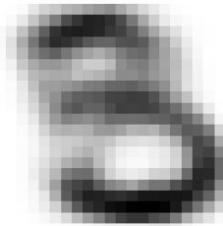
(0.70, -1.06)



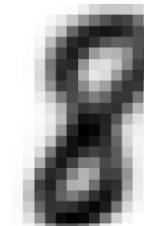
(1.36, -2.46)



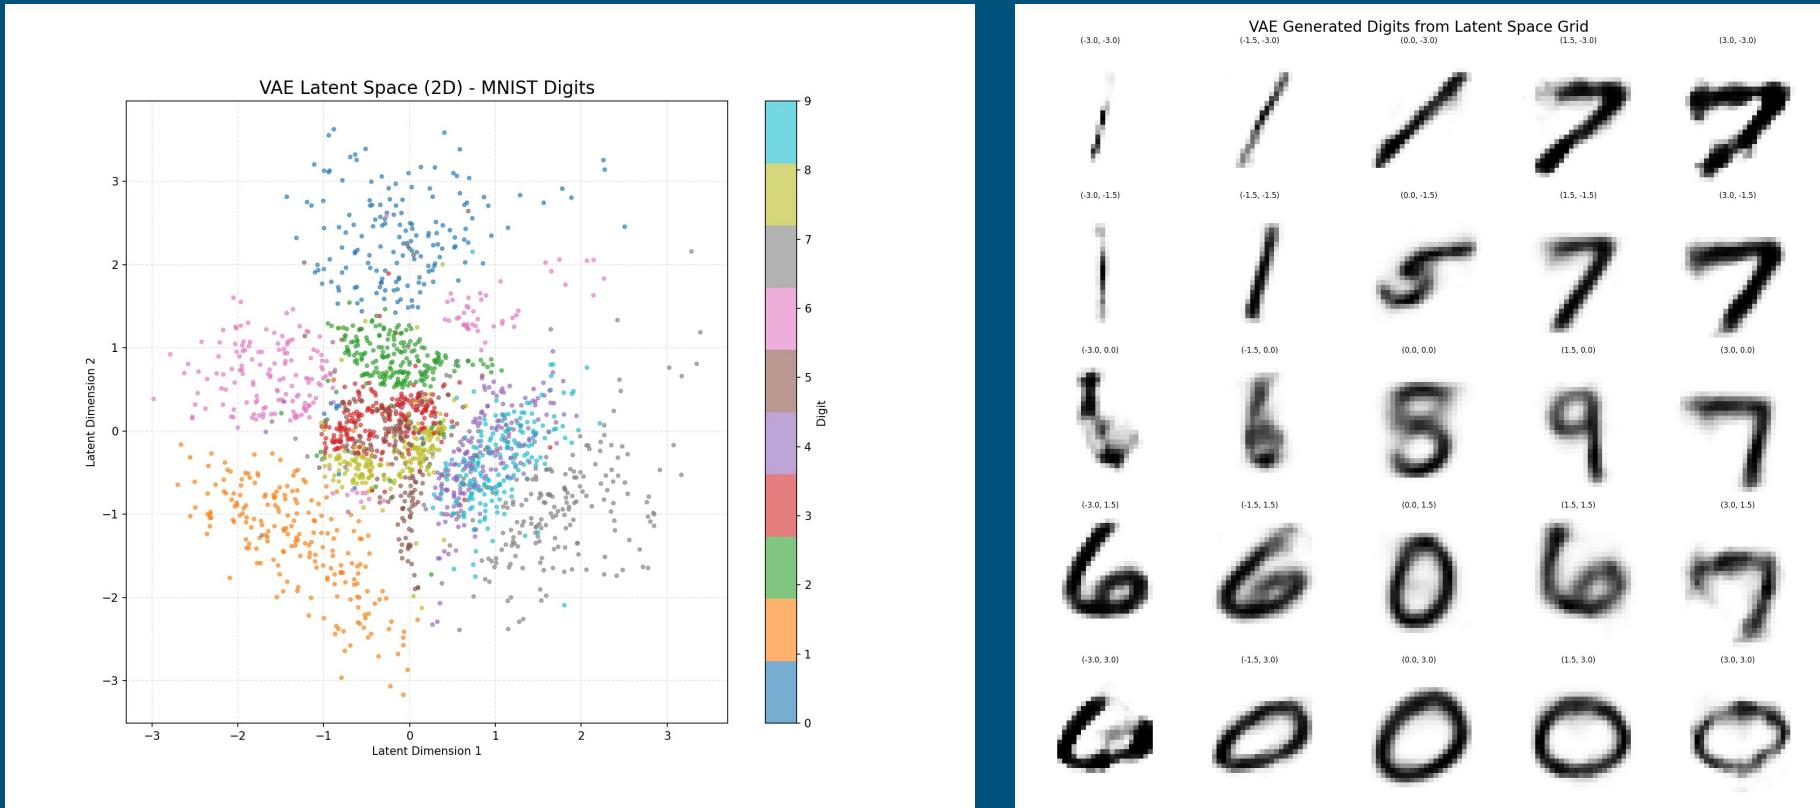
(0.15, 0.47)



(-0.67, -0.46)



# Espacio Latente ( $E = 0.175$ )



# Si modificamos la arquitectura....

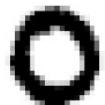
---

- ❑ Arquitectura:
    - ❑ Encoder: 784 -> 256 -> 128
    - ❑ Espacio Latente: 2
    - ❑ Decoder: 128 -> 256 -> 784
  - ❑ Optimizador: ADAM
  - ❑ Learning Rate: 0.001
- 
- ❑ Funciones de activación:
    - ❑ ReLU (Capas ocultas)
    - ❑ Sigmoidea (Última capa)
  - ❑ Epochs: 200
  - ❑ Batch Size: 64 (Mini-batch)

# Reconstrucción ( $E = 0.170$ )

VAE Reconstruction: Original vs Reconstructed (One per Digit)

Original 0



Original 1



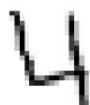
Original 2



Original 3



Original 4



Original 5



Original 6



Original 7



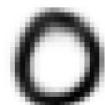
Original 8



Original 9



Reconstructed



Reconstructed



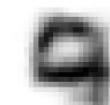
Reconstructed



Reconstructed



Reconstructed



Reconstructed



Reconstructed



Reconstructed



Reconstructed



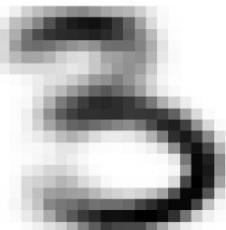
Reconstructed



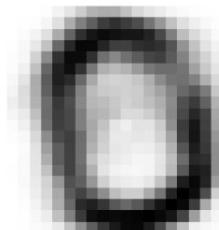
# Generación ( $E = 0.170$ )

VAE Random Generation with Latent Coordinates

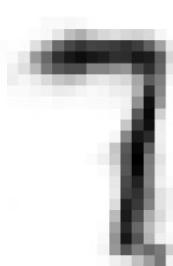
(-0.20, 2.23)



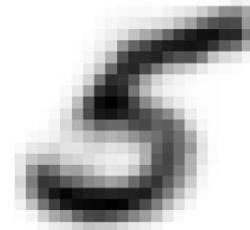
(-0.86, 1.62)



(1.98, -1.64)



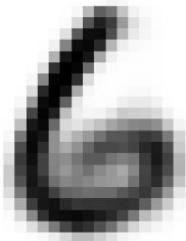
(0.01, 0.22)



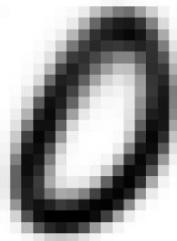
(-1.34, 0.52)



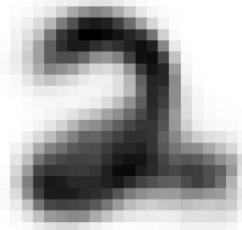
(-0.58, -0.39)



(-1.01, 0.93)



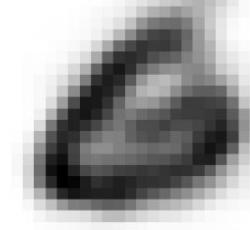
(-1.30, -0.03)



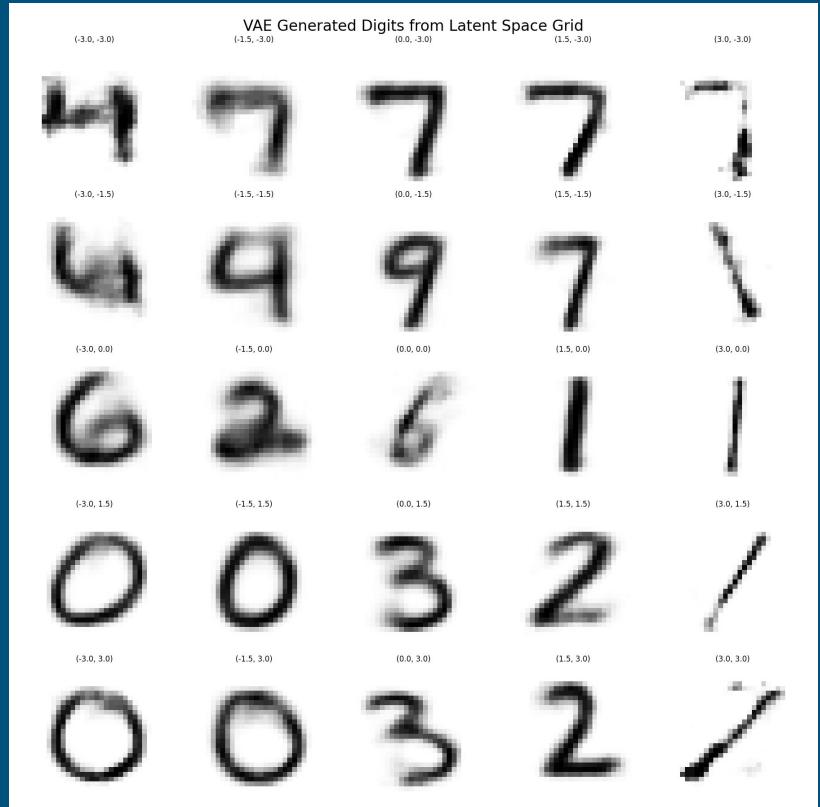
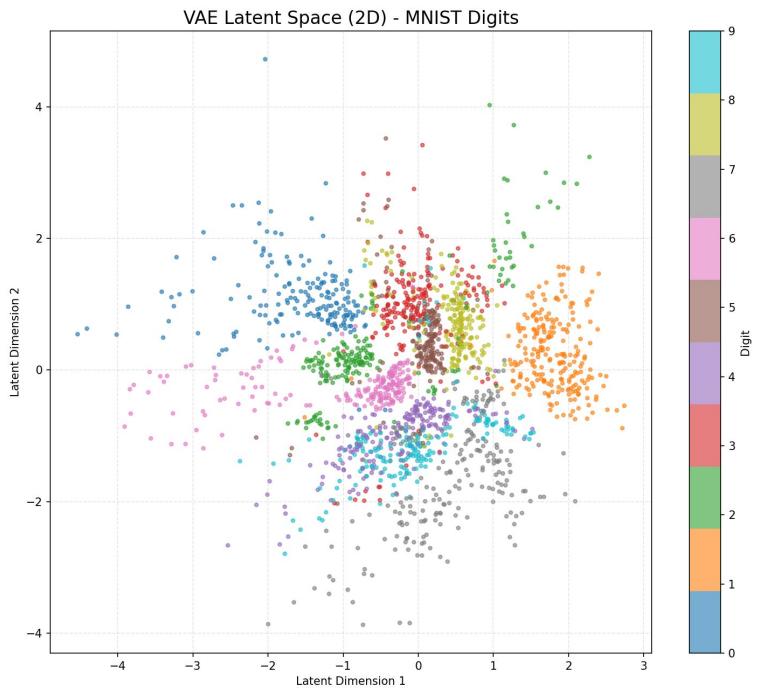
(-0.69, 0.77)



(-1.74, 0.27)



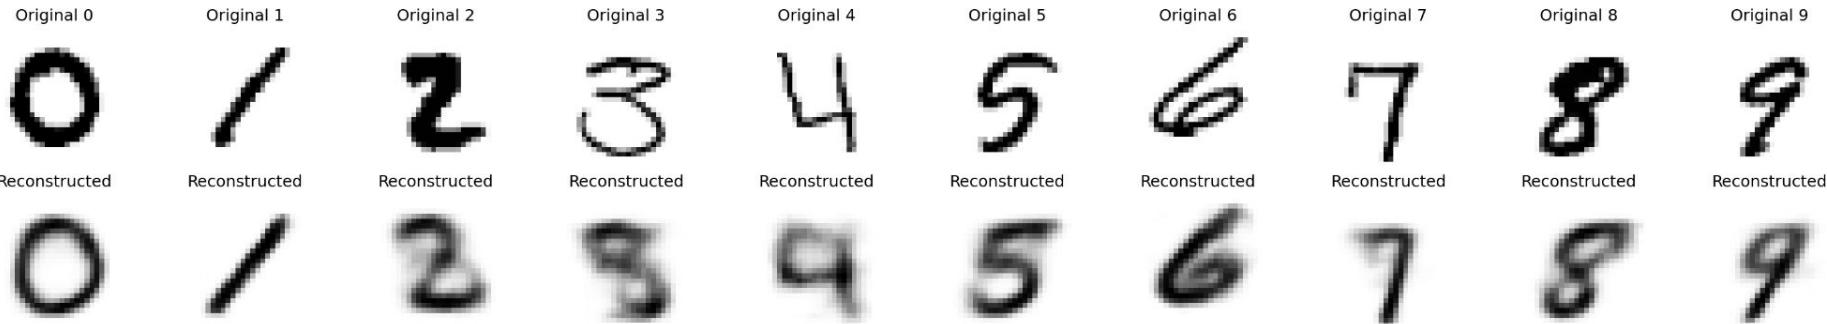
# Espacio Latente ( $E = 0.170$ )



# Comparación

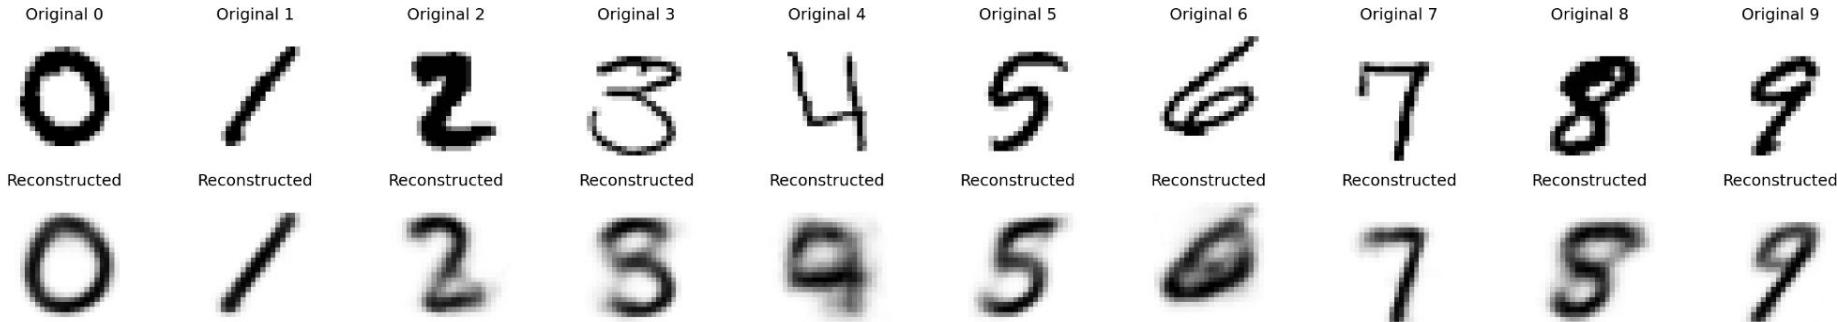
# Reconstrucción 1 ( $E = 0.175$ )

VAE Reconstruction: Original vs Reconstructed (One per Digit)



# Reconstrucción 2 ( $E = 0.170$ )

VAE Reconstruction: Original vs Reconstructed (One per Digit)



# Conclusiones

---

- El DAE necesita un mínimo de copias con ruido en el dataset para aprender a corregir errores y memorizar patrones base. Aún así es muy difícil arreglar errores producto de patrones similares.
- En el VAE, la arquitectura  $128 \rightarrow 64$  demostró ser suficiente para modelar dígitos MNIST, con una mejora marginal al duplicar parámetros ( $256 \rightarrow 128$ : 0.170 vs 0.175)
- El VAE logró separar automáticamente las 10 clases de dígitos en clusters diferenciados en un espacio 2D, sin supervisión de etiquetas durante el entrenamiento